


Same question, seven doors

What AI research tools agree on, where they part ways, and where your own judgment has to stay in the loop.

1 • THE RESEARCH ARC, COLORED BY HOW MUCH THE TYPICAL AI MOVE DISPLACES YOUR REASONING




2 • ONE QUESTION, SEVEN DOORS

“Does using generative AI during learning or research tasks reduce people's later independent performance, such as recall, understanding, or skill?”

ENDPOINT	PAPERS	ONLY HERE
PubMed	20	19
Consensus	20	16
Zotero	12	10
alphaXiv	10	9
Elicit	...	-
ResearchRabbit	...	-
bioRxiv / medRxiv	n/a	-

2.1 source overlap (mean pairwise Jaccard $\times 100$ of the top-20 lists)

3 • THE CONVERGENCE MAP, ALL QUESTIONS



4 • THREE CLAIMS, CALIBRATED TO THEIR EVIDENCE

“Individually better, collectively narrower”

Doshi & Hauser 2024 · Online experiment, short stories

“The clearest field evidence that performance and learning come apart under AI”

Bastani et al. 2025 · Randomized field trial, high-school mathematics

“The feeling of discovery is not evidence”

Laukkonen et al. 2020 · Preregistered experiment

5 • THE RULE

Generate, then check: predict, retrieve, or draft first, then consult the source or the AI.

Write your frame, prediction, or 150-word synthesis before you open the tool. Then ask it what you missed.

BUDGET-AND-STOP CARD

Value	What decision does this inform, and what is a good answer worth?
Budget	How many hours have similar questions taken me before?
Pressure by phase	Loose for Learn or Discover; tight for Define, Investigate, Rapid Answer.
Minimum targets	What result, by what date, would tell me this approach has failed?
Re-bet date	On what date do I decide deliberately whether to continue?
Opportunity cost	What am I not doing while I do this?

